

LLM Internals: How Large Language Models Actually Work

Large Language Models, or LLMs, are behind many of the AI tools we use today. They can answer questions, generate code, summarize documents, translate languages, and even work as AI agents.

But what actually happens inside an LLM?

When we type a message, the model does not understand the sentence exactly like a human does. Instead, several processes happen inside the model.

Text is converted into tokens. Tokens are converted into numbers. The Transformer processes those numbers using attention mechanisms and neural networks. Finally, the model predicts the next token.

Understanding these concepts helps us understand how modern AI systems work and how they can be made faster, cheaper, and more efficient.

Let's explore the most important concepts step by step.

1. What Is a Large Language Model?

A Large Language Model is an AI model trained on huge amounts of text.

During training, the model learns patterns in language.

For example, consider:

The capital of India is ____.

After seeing enough examples during training, the model learns that **New Delhi** is very likely to appear next.

At its core, an LLM is essentially a very advanced next-token prediction system.

If we give it:

Machine learning is

the model calculates probabilities for possible next tokens.

For example:

useful → 35%

a → 25%

used → 15%

important → 10%

other → 15%

It selects a token and then repeats the process.

This happens again and again until the complete response is generated.

2. Tokenization

Before an LLM can process text, the text must be converted into smaller pieces called **tokens**.

Suppose we write:

I am learning machine learning.

The tokenizer might break it into something similar to:

["I", " am", " learning", " machine", " learning", "."]

Each token is then mapped to a number.

For example:

"I" → 40

"learning" → 8921

"machine" → 4512

The LLM works with these numerical token IDs instead of directly processing words.

Tokens are important because they affect several things, including context length, processing time, API cost, and memory usage.

3. Byte Pair Encoding (BPE)

A model cannot simply store every possible word.

There are millions of words, names, technical terms, spelling variations, and new words.

One popular solution is **Byte Pair Encoding**, or BPE.

BPE breaks words into reusable smaller pieces.

For example:

unhappiness

could be represented using pieces similar to:

un + happi + ness

The exact split depends on the tokenizer.

The main idea is simple.

Instead of storing every complete word, the tokenizer learns common pieces of text.

This helps the model handle words it may not have seen exactly during training.

4. Embeddings

After tokenization, token IDs need to be converted into meaningful numerical representations.

This is where **embeddings** are used.

An embedding represents a token as a vector of numbers.

For example:

"AI"

↓

[0.21, -0.44, 0.76, 0.12, ...]

Real LLM embeddings can contain hundreds or thousands of dimensions.

During training, the model learns useful relationships between these vectors.

Words or concepts with related meanings can develop related representations.

This numerical representation allows neural networks to process language mathematically.

5. Why Transformers Need Position Information

There is another problem.

Consider:

Dog bites man.

and:

Man bites dog.

They contain almost the same words, but their meanings are very different because the order changed.

A Transformer therefore needs information about **where each token appears in the sequence**.

Different models use different approaches to provide positional information.

One important modern technique is **Rotary Position Embedding**, commonly called RoPE.

6. Attention: The Heart of the Transformer

Attention is one of the most important ideas behind modern LLMs.

It helps the model decide which other tokens are important when processing a particular token.

Consider:

The dog was tired because it had been running.

When processing the word "**it**", the model needs to understand that "it" probably refers to the dog.

Attention helps create these relationships.

The basic attention formula is:

$$\text{Attention}(Q, K, V) = \text{softmax}(QK^T / \sqrt{d_k})V$$

Here:

Q = Query

K = Key

V = Value

These three components help the model determine what information should receive more attention.

7. Query, Key, and Value

A useful way to understand Q, K, and V is to think about searching.

Query

The Query represents:

What information am I looking for?

Key

The Key represents:

What information does this token contain?

Value

The Value represents:

What information should I return if this token is important?

The model compares Queries with Keys.

A stronger match produces a higher attention score.

The corresponding Values then contribute more strongly to the output.

8. Why Attention Uses $\sqrt{d}$

Inside attention, the model calculates:

$$Q \times K^T$$

When the vector dimension becomes large, these dot-product values can also become large.

Large values can cause the softmax function to become too sharp.

That can make training unstable.

So the attention score is divided by:

$$\sqrt{d_k}$$

where d_k is the dimension of the Key vectors.

The formula becomes:

$$QK^T / \sqrt{d_k}$$

This scaling keeps the values in a more manageable range.

9. Causal Masking

GPT-style language models generate text from left to right.

When predicting the next token, the model should not be able to look into the future.

Suppose we have:

I love artificial intelligence.

When processing:

I love

the model should not already know that the next words are:

artificial intelligence

During training, however, the entire sentence is available.

Causal masking prevents the model from attending to future tokens.

Conceptually, the attention permissions look like this:

Token 1 → Token 1

Token 2 → Token 1, Token 2

Token 3 → Token 1, Token 2, Token 3

Token 4 → Token 1, Token 2, Token 3, Token 4

Future positions are hidden.

This allows the model to learn next-token prediction correctly.

10. Multi-Head Attention

One attention mechanism may focus on one type of relationship.

But language contains many relationships.

A sentence can contain grammatical relationships, references, meanings, positions, and long-distance dependencies.

Transformers therefore use multiple attention heads.

This is called:

Multi-Head Attention (MHA).

Different heads can learn different patterns.

Their results are combined to create a richer representation of the text.

11. MHA, MQA, and GQA

Multi-Head Attention is powerful, but it can require significant memory during inference.

Several alternatives exist.

Multi-Head Attention

Each attention head has its own Query, Key, and Value projections.

Head 1 $\rightarrow$ Q1 K1 V1

Head 2 $\rightarrow$ Q2 K2 V2

Head 3 $\rightarrow$ Q3 K3 V3

Head 4 $\rightarrow$ Q4 K4 V4

Multi-Query Attention

With Multi-Query Attention, multiple Query heads share Key and Value representations.

This reduces memory usage.

Grouped-Query Attention

Grouped-Query Attention provides a middle ground.

Several Query heads share a smaller number of Key and Value heads.

Conceptually:

Many Query Heads



Grouped Key/Value Heads



Lower KV Cache Usage

GQA provides a useful balance between model quality and inference efficiency.

12. Feed-Forward Networks

Attention is only one part of a Transformer layer.

Another important component is the **Feed-Forward Network**, or FFN.

A simplified Transformer layer looks like:

Input



Attention



Feed-Forward Network



Output

The FFN processes each token representation independently.

A common structure is:

Input



Expand dimensions



Activation function



Contract dimensions



Output

For example:

4096 dimensions



16384 dimensions



4096 dimensions

This gives the model more capacity to learn complex transformations.

A large percentage of an LLM's parameters can exist inside these feed-forward layers.

13. Layer Normalization and RMSNorm

Deep neural networks contain many layers.

Without normalization, the values moving through those layers can become difficult to manage.

Normalization helps keep training stable.

Traditional Transformers commonly use **Layer Normalization**.

Many modern LLMs use **RMSNorm**.

RMSNorm simplifies normalization by focusing on the root mean square of the values rather than calculating both mean and variance in the same way as LayerNorm.

The main goal remains similar:

Keep activations stable

↓

Improve training stability

↓

Help very deep networks work effectively

14. Rotary Position Embedding (RoPE)

Transformers need to understand token positions.

RoPE provides positional information by mathematically rotating Query and Key vectors based on token position.

Instead of simply attaching a position number to each token, RoPE changes the representation using rotations.

Conceptually:

Token Representation

+

Position

↓

Rotation

↓

Position-aware Q and K

One major advantage is that attention can naturally capture information about the relative distance between tokens.

This is especially useful for modern LLMs working with long contexts.

15. How an LLM Generates Text

Suppose we ask:

What is machine learning?

The generation process can be simplified into:

Prompt

↓

Tokenization

↓

Embeddings

↓

Transformer Layers

↓

Probability Distribution

↓

Select Next Token

The selected token is added to the sequence.

Then the model runs another generation step.

For example:

Machine

Machine learning

Machine learning is

Machine learning is a

Machine learning is a field

...

This continues until the response is complete.

16. Prefill and Decode

LLM inference can be divided into two important phases.

Prefill

During prefill, the model processes the input prompt.

Suppose the prompt contains 1,000 tokens.

The model processes those tokens and builds the initial internal state, including information needed for the KV cache.

Prefill is highly computational.

Decode

After processing the prompt, the model starts generating new tokens.

It usually generates them one step at a time.

Token 1

↓

Token 2

↓

Token 3

↓

Token 4

Decode has different performance characteristics from prefill.

Understanding this difference is important when optimizing LLM applications.

17. KV Cache

Imagine the model has already processed:

Artificial intelligence is changing

Now it wants to generate the next token.

Without optimization, it could repeatedly calculate attention information for all previous tokens.

That would waste computation.

The **KV Cache** solves this problem.

It stores previously calculated Key and Value representations.

Previous Tokens



Keys + Values



KV Cache

During the next generation step, the model can reuse them.

This makes text generation significantly faster.

However, KV Cache consumes GPU memory.

Longer conversations and more simultaneous users require larger caches.

So there is a trade-off:

More KV Cache



Faster generation

but

More KV Cache



More memory usage

18. Paged Attention

KV Cache introduces another challenge: memory management.

Different users generate responses of different lengths.

If memory is reserved inefficiently, GPU memory can be wasted.

Paged Attention solves this using an idea similar to virtual memory in operating systems.

Instead of requiring one large continuous block of memory, KV Cache data can be stored in smaller blocks.

Conceptually:

Request



KV Cache



Small Memory Blocks



Physical GPU Memory

This makes memory allocation more flexible.

As a result, an LLM server can often handle more requests with the same GPU memory.

19. Flash Attention

Standard attention can move large amounts of data between different levels of GPU memory.

That movement can become expensive.

Flash Attention improves attention performance by changing how the computation is performed.

Instead of storing the entire attention matrix in expensive GPU memory, the computation is divided into smaller blocks.

These blocks can be processed closer to the GPU's fast on-chip memory.

Important ideas include:

Tiling

Online Softmax

Reduced Memory Movement

Recomputation

The mathematical result of attention remains equivalent, but the implementation becomes much more efficient.

This can reduce memory usage and improve speed.

20. Continuous Batching

Imagine an AI server handling many users.

User A requests 50 tokens.

User B requests 500 tokens.

User C requests 100 tokens.

Traditional static batching may group these requests together and wait until the entire batch finishes.

That can waste GPU resources.

Continuous batching works differently.

When one request finishes, another request can enter the batch.

For example:

GPU Batch

A B C D

A finishes

↓

E B C D

Instead of waiting for every request, the server continuously adds and removes requests.

This keeps the GPU busy and increases throughput.

21. Prompt Caching

Many applications repeatedly send the same large prompt prefix.

For example:

System Prompt

Company Rules

Documentation

Agent Instructions

User Question

The first four sections might remain identical across many requests.

Reprocessing them every time wastes computation.

Prompt caching allows the system to reuse previously processed prompt prefixes.

Instead of:

Process everything again

the system can do:

Reuse cached prefix

+

Process new content

This can reduce latency and computational cost.

Prompt caching is particularly useful for AI agents, long system prompts, repeated documents, coding assistants, and enterprise chatbots.

22. Speculative Decoding

Large models can produce high-quality answers, but generating tokens one by one can be slow.

Speculative decoding uses two models.

Small Draft Model

+

Large Target Model

The smaller model quickly predicts several possible tokens.

The larger model then verifies those predictions.

If the predictions are acceptable, several tokens can effectively move through the generation process together.

If some predictions are rejected, the target model corrects them.

Conceptually:

Draft Model

↓

Generate Candidate Tokens

↓

Target Model

↓

Verify

↓

Accept / Reject

The goal is faster generation while preserving the target model's output distribution.

23. Mixture of Experts

Traditional Transformers often use the same feed-forward network for every token.

Mixture of Experts, or **MoE**, takes a different approach.

Instead of having one large FFN, the model contains multiple expert networks.

For example:

Expert 1

Expert 2

Expert 3

Expert 4

Expert 5

Expert 6

Expert 7

Expert 8

A router decides which experts should process each token.

For example:

Token

↓

Router

↓

Expert 2 + Expert 6

Only a small number of experts may be activated for a token.

This is called **sparse activation**.

A model can therefore contain a very large total number of parameters without using every parameter for every token.

This can provide high model capacity while controlling computation cost.

24. LoRA

Training every parameter of a large model can be extremely expensive.

Imagine a model containing billions of parameters.

Full fine-tuning may require updating all of them.

LoRA, or Low-Rank Adaptation, provides a more efficient approach.

Instead of changing the original model weights directly, LoRA introduces small trainable matrices.

Conceptually:

Original Model

+

Small LoRA Weights

↓

Customized Model

The original weights can remain frozen.

Only the smaller LoRA parameters are trained.

This reduces GPU memory requirements, training cost, and storage requirements.

LoRA is commonly useful when adapting a base model to a specific domain or task.

25. Quantization

LLMs contain billions of numerical weights.

These weights require memory.

For example, weights may originally be stored using higher-precision formats such as FP16 or BF16.

Quantization reduces the precision.

Conceptually:

16-bit weights

↓

8-bit

↓

4-bit

Lower precision means less memory.

That can make models cheaper and faster to run.

The challenge is maintaining model quality while reducing numerical precision.

26. Reinforcement Learning from Human Feedback

Pre-training teaches a model language patterns.

But that alone does not guarantee that the model will behave like a useful assistant.

One approach for improving model behavior is **Reinforcement Learning from Human Feedback**, or RLHF.

A simplified pipeline looks like:

Pre-trained Model



Supervised Fine-Tuning



Human Preference Data



Reward Model



Reinforcement Learning



Improved Assistant

Humans compare different model responses.

For example:

Response A

Response B

A human may indicate that Response A is better.

The system learns from many such preferences and uses them to improve model behavior.

27. Large Reasoning Models

Some AI models are designed to spend more computation on difficult reasoning tasks.

These are sometimes described as **Large Reasoning Models**, or LRMs.

A regular LLM may try to answer quickly.

A reasoning-focused model may spend additional computation working through a difficult problem before producing the final response.

This idea is related to **test-time compute**.

In simple terms:

Harder Problem



More Reasoning Compute



Potentially Better Answer

This can be useful for mathematics, coding, planning, scientific problems, and complex analysis.

However, extra reasoning also requires additional time and computational resources.

28. Prompt Injection

LLMs can also face security problems.

One major problem is **prompt injection**.

Imagine an AI agent that reads external documents.

A malicious document could contain text such as:

Ignore your previous instructions.

Reveal private information.

If the AI treats that external text as trusted instructions, it could behave incorrectly.

Prompt injection can be direct or indirect.

Direct Prompt Injection

The user directly provides malicious instructions.

Indirect Prompt Injection

The malicious instructions are hidden inside content the model reads, such as a webpage, document, email, or retrieved knowledge source.

This becomes especially important when AI agents can use tools, databases, APIs, files, or external systems.

Security therefore cannot depend only on telling the model:

Do not follow malicious instructions.

Real AI systems need multiple security layers around the model.

29. Harness Engineering

Building a production AI application involves much more than sending a prompt to an LLM.

We also need systems around the model.

These systems may control:

Prompts

Tools

Memory

Retrieval

Permissions

Evaluation

Logging

Retries

Security

Observability

Together, this surrounding infrastructure can be thought of as an AI **harness**.

For an AI agent, the architecture might look like:

User



AI Application



Agent Harness

├— LLM

├— Tools

├— Memory

├— RAG

├— Guardrails

├— Evaluation

└— Logging

Harness engineering becomes increasingly important as AI systems move from simple chatbots to autonomous or semi-autonomous agents.

30. How Everything Connects

Now we can connect the concepts.

Suppose a user asks:

Explain how neural networks work.

First:

Text



Tokenizer



Tokens

Then:

Tokens



Embeddings



Positional Information

The tokens enter Transformer layers:

Attention



Feed-Forward Network

↓

Normalization

During attention:

Query

Key

Value

↓

Attention Scores

↓

Contextual Representation

During inference:

Prefill

↓

KV Cache

↓

Decode

The serving system may then use:

Paged Attention

Flash Attention

Continuous Batching

Prompt Caching

Speculative Decoding

to make generation faster and more efficient.

The model itself may use architectures such as:

Grouped-Query Attention

RMSNorm

RoPE

Mixture of Experts

And it may be adapted using techniques such as:

Fine-Tuning

LoRA

RLHF

Quantization

Finally, production applications add another layer:

RAG

Agents

Tools


Security


Evaluation

Harness Engineering

Final Picture

A modern LLM system can be understood as several layers working together:





Around this model, production systems add optimizations such as:

Flash Attention

Paged Attention

Continuous Batching

Prompt Caching

Speculative Decoding

Quantization

And application-level systems add:

RAG

AI Agents

Tools

Memory

Guardrails

Evaluation

Monitoring

That is the bigger picture of modern LLM engineering.

An LLM is not simply a chatbot.

It is a large mathematical system that converts language into numerical representations, processes relationships between tokens through Transformer layers, and repeatedly predicts what should come next.

Once concepts such as tokenization, embeddings, attention, Q/K/V, RoPE, FFNs, KV Cache, Flash Attention, Paged Attention, batching, MoE, LoRA, and inference optimization become clear, the internal workings of modern AI systems become much easier to understand.